

Resilient distributed datasets in PySpark


INTRODUCTION TO PYSPARK



Benjamin Schmidt
Data Engineer

What is parallelization in PySpark?

- Automatically parallelizing data and computations across multiple nodes in a cluster
- Distributed processing of large datasets across multiple nodes
- Worker nodes process data in parallel, combining at the end of the task
- Faster processing at scale (think gigabytes or even terabytes)



Understanding RDDs

RDDs or *Resilient Distributed Datasets*:

- Distributed data collections across a cluster with automatic recovery from node failures
- Good for large scale data
- Immutable and can be transformed using operations like `map()` or `filter()`, with actions like `collect()` or `parallelize()` to retrieve results or create RDDs

Creating an RDD

```
# Initialize a Spark session
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("RDDExample").getOrCreate()

# Create a DataFrame from a csv
census_df = spark.read.csv("/census.csv")
# Convert DataFrame to RDD
census_rdd = census_df.rdd

# Show the RDD's contents using collect()
census_rdd.collect()
```

Showing Collect

```
# Collect the entire DataFrame into a local Python list of Row objects
data_collected = df.collect()

# Print the collected data
for row in data_collected:
    print(row)
...
```

RDDs vs DataFrames

DataFrames

- High-level: Optimized for ease of use
- SQL Like Operations: Work with SQL-like queries and perform complex operations with less code
- Schema Information: Contain Columns and types like an SQL Table

RDDs

- Low-level: More flexible but requiring more lines of code for complex operations
- Type Safety: Preserve data types but don't have the optimization benefits of DataFrames
- No Schema: Harder to work with structured data like SQL or relational data
- Large Scaling
- Very very verbose compared to DataFrames and poor at analytics

Some useful functions and methods

- `map()` : method applies functions (including ones we write like a lambda function) across a dataset like: `rdd.map(map_function)`
- `collect()` : collects data from across the cluster like: `rdd.collect()`

Let's practice!

INTRODUCTION TO PYSPARK

Intro to Spark SQL

INTRODUCTION TO PYSPARK



Benjamin Schmidt
Data Engineer

What is Spark SQL

- Module in Apache Spark for structured data processing
- Allows us to run SQL queries alongside data processing tasks
- Seamless combination of Python and SQL in one application
- DataFrame Interfacing: Provides programmatic access to structured data

Creating temp tables

```
# Initialize Spark session
spark = SparkSession.builder.appName("Spark SQL Example").getOrCreate()

# Sample DataFrame
data = [("Alice", "HR", 30), ("Bob", "IT", 40), ("Cathy", "HR", 28)]
columns = ["Name", "Department", "Age"]
df = spark.createDataFrame(data, schema=columns)

# Register DataFrame as a temporary view
df.createOrReplaceTempView("people")

# Query using SQL
result = spark.sql("SELECT Name, Age FROM people WHERE Age > 30")
result.show()
```

Deeper into temp views

- Temp Views protect the underlying data while doing analytics
- Loading from a CSV uses methods we already know

```
df = spark.read.csv("path/to/your/file.csv", header=True, inferSchema=True)
```

```
# Register DataFrame as a temporary view  
df.createOrReplaceTempView("employees")
```

Combining SQL and DataFrame operations

```
# SQL query result
query_result = spark.sql("SELECT Name, Salary FROM employees WHERE Salary > 3000")
# DataFrame transformation
high_earners = query_result.withColumn("Bonus", query_result.Salary * 0.1)
high_earners.show()
```

Let's practice!

INTRODUCTION TO PYSPARK

PySpark aggregations

INTRODUCTION TO PYSPARK



Benjamin Schmidt
Data Engineer

PySpark SQL aggregations overview

- Common SQL aggregations work with `spark.sql()`

```
# SQL aggregation query
spark.sql("""
    SELECT Department, SUM(Salary) AS Total_Salary, AVG(Salary) AS Average
    FROM employees
    GROUP BY Department
""").show()
```


Combining DataFrame and SQL operations

```
# Filter salaries over 3000
filtered_df = df.filter(df.Salary > 3000)

# Register filtered DataFrame as a view
filtered_df.createOrReplaceTempView("filtered_employees")

# Aggregate using SQL on the filtered view
spark.sql("""
    SELECT Department, COUNT(*) AS Employee_Count
    FROM filtered_employees
    GROUP BY Department
""").show()
```

Handling data types in aggregations

```
# Example of type casting
data = [("HR", "3000"), ("IT", "4000"), ("Finance", "3500")]
columns = ["Department", "Salary"]
df = spark.createDataFrame(data, schema=columns)

# Convert Salary column to integer
df = df.withColumn("Salary", df["Salary"].cast("int"))

# Perform aggregation
df.groupBy("Department").sum("Salary").show()
```

RDDs for aggregations

```
# Example of aggregation with RDDs
rdd = df.rdd.map(lambda row: (row["Department"], row["Salary"]))

rdd_aggregated = rdd.reduceByKey(lambda x, y: x + y)

print(rdd_aggregated.collect())
```

Best practices for PySpark aggregations

- Filter early: Reduce data size before performing aggregations
- Handle data types: Ensure data is clean and correctly typed
- Avoid operations that use the entire dataset: Minimize operations like `groupBy()`
- Choose the right interface: Prefer DataFrames for most tasks due to their optimizations
- Monitor performance: Use `explain()` to inspect the execution plan and optimize accordingly

Key takeaways

- PySpark SQL Aggregations: Functions like `SUM()` and `AVERAGE()` for summarizing data
- DataFrames and SQL: Combining both approaches for flexible data manipulation
- Handling Data Types: Addressing issues with type mismatches during aggregations
- RDDs vs DataFrames: Understanding the trade-offs and choosing the right tool

Let's practice!

INTRODUCTION TO PYSPARK

PySpark at scale

INTRODUCTION TO PYSPARK



Benjamin Schmidt
Data Engineer

Leveraging scale

- Pyspark works effectively with gigabytes and terabytes of data
- Using PySpark, speed and efficient processing is the goal
- Understanding PySpark execution gets even more efficiencies
- Use broadcast to manage the whole cluster

```
joined_df = large_df.join(broadcast(small_df),  
                           on="key_column", how="inner")  
  
joined_df.show()
```


Execution plans

```
# Using explain() to view the execution plan
df.filter(df.Age > 40).select("Name").explain()
```

```
== Physical Plan ==
*(1) Filter (isNotNull(Age) AND (Age > 30))
+- Scan ExistingRDD[Name:String, Age:Int]
```

Caching and persisting DataFrames

- Caching: Stores data in memory, for faster access for smaller datasets
- Persisting: Stores data in different storage levels for larger datasets

```
df = spark.read.csv("large_dataset.csv", header=True, inferSchema=True)
```

```
# Cache the DataFrame
```

```
df.cache()
```

```
# Perform multiple operations on the cached DataFrame
```

```
df.filter(df["column1"] > 50).show()
```

```
df.groupBy("column2").count().show()
```

Persisting DataFrames with different storage levels

```
# Persist the DataFrame with storage level
from pyspark import StorageLevel

df.persist(StorageLevel.MEMORY_AND_DISK)

# Perform transformations
result = df.groupBy("column3").agg({"column4": "sum"})
result.show()

# Unpersist after use
df.unpersist()
```

Optimizing PySpark

- Small Subsections: The more data that gets used, the slower the operation: Pick tools like `map()` over `groupby()` due to selectivity of methods
- Broadcast Joins: Broadcast will use all compute, even on smaller datasets
- Avoid Repeated Actions: Repeated actions on the same data costs time and compute, without any benefit

Let's practice!

INTRODUCTION TO PYSPARK

What have we learned?

INTRODUCTION TO PYSPARK



Benjamin Schmidt
Data Engineer

What you did

- Learned about PySparks clusters
- PySpark critical syntax
- RDDs and DataFrames
- Spark SQL



What you haven't done (yet)

- Cluster management
- Complex job optimization
- PySpark at scale
- Machine learning

What you can do next on DataCamp

- Big Data Fundamentals with PySpark
- Cleaning Data with PySpark
- Machine Learning with PySpark

**Keep going and
practicing**

INTRODUCTION TO PYSPARK